

CUDA: understand to be a genuine user

takigawa

May 20, 2025

The university of Tokyo, EEIC, Taura Lab

1. What is CUDA?: Introduction
2. A CUDA program for beginners
3. How CUDA works
4. Optimize CUDA program
5. Practice Problems

What is CUDA?

1. What is CUDA?: Introduction
2. A CUDA program for beginners
3. How CUDA works
4. Optimize CUDA program
5. Practice Problems

CUDA is the abstraction of GPU(s) for programmers

GPU (Graphical Processing Unit) is a device separated from CPU (Central PU).

- Code that runs on GPU must be designated as a kernel
- Data must be copied between CPU and GPU(s)
- A GPU is often called a *"device"*
- A CPU is often called a *"host"*



Figure 1: Host

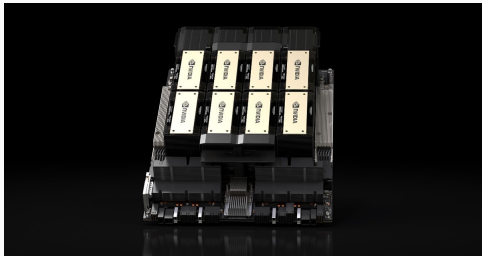


Figure 2: Device

CUDA is the abstraction of GPU(s) for programmers

CUDA is a platform for parallel computing on NVIDIA GPU.

- language extension: C/C++, Fortran
- tools: compiler(nvcc), debugger(cuda-gdb), profiler(Nsight Systems)
- APIs: Driver API, Runtime API
- Libraries: cuBLAS, cuFFT, cuDNN, etc.

Their common goal is to provide programmers with a "good" abstraction of GPU(s).
How "good"? : usable, simple, highly affine to hardware(=easy to bring out the performance)

To compile/run CUDA programs: NVCC

```
1      nvcc -arch=sm_90 programs.cu%
```

- compile with `nvcc` command
- the natural extension of CUDA program is `.cu`
- `-arch` flag designates GPU Architecture
 - `compute_XX`: Virtual architecture
 - `sm_XX`: Physical architecture (SM generation)

Interlude: How to know the proper architecture

Use `cudaGetDeviceXXXXX` APIs and get device query.

`device_query.cu`

```
1 #include <cuda_runtime.h>
2 int deviceCount;
3 cudaError_t error = cudaGetDeviceCount(&deviceCount);
4 cudaDeviceProp deviceProp;
5 cudaGetDeviceProperties(&deviceProp, i);
6
7 printf("Device %d: %s\n", i, deviceProp.name);
8 printf("  Compute capability: %d.%d\n", deviceProp.major, deviceProp.minor);
9 printf("  Total global memory: %.2f GB\n", deviceProp.totalGlobalMem / (1024.0
    * 1024.0 * 1024.0));
```

or

```
1  nvidia-smi --query-gpu=compute_cap
```

The architecture for GH200 is `sm_90`.

A CUDA program for beginners

1. What is CUDA?: Introduction
2. A CUDA program for beginners
 - 2.1 Kernels; writing and launching
 - 2.2 Host-Device Data Communication (+ Synchronization)
3. How CUDA works
4. Optimize CUDA program
5. Practice Problems

Setting environment: Using GH200(s) on miyabi

1. <https://miyabi-www.jcahpc.jp/login> にアクセスし、パスワード初期化を選択する
2. 指示にしたがい、Miyabi 利用支援ポータルにアクセスする
3. ドキュメント閲覧/Miyabi システム利用手引書 をダウンロードする (Strongly recommended)
4. 手引書の P.9 システム初回ログイン時の設定 の手順を完了する
5. 手引書の P.22 SSH ログイン/初回ログイン の手順を完了する (必ず緊急用スクラッチコードを控えること)
6. (必要に応じて) エディタから 2 回目ログインを行う

2 要素認証が必須である。

【注意】 ログインノード `/home/cXXXXX` ではなく、計算ノード `/work/gc64/cXXXXX` で作業する

Sample program #1: hello_world.cu

See <https://github.com/gunnersgoestocl/cuda-introduction/tree/main/tutorial-legacy> for more information.

```
1 #include <stdio.h>
2 #include <cuda_runtime.h>           // for Runtime APIs
3
4 __global__ void hello(){           // kernel function
5     printf("Hello CUDA World !!\n");
6 }
7
8 int main() {
9     hello<<< 2, 4 >>>();           // launch kernel
10    cudaDeviceSynchronize();        // wait until kernel completes
11    return 0;
12 }
13
```

3 files are required for execution on miyabi

- .cu file: CUDA user program
- makefile: compile and clean
- shell script: to submit batch job, see official docs for more info

```
1 NVCC := nvcc
2 NVCCFLAGS := -arch=sm_90 -O3
3 # .cu ファイルから実行ファイルを生成
4 CUDA_EXECUTABLES := $(patsubst %.cu,%,$(wildcard
    *.cu))
5 # デフォルトのターゲット
6 all: $(C_EXECUTABLES) $(CUDA_EXECUTABLES)
7 # .cu ファイルから実行ファイルを生成
8 %: %.cu
9 ^^I$(NVCC) $(NVCCFLAGS) $< -o $@
10 # clean ターゲットの定義
11 .PHONY: clean
12 clean:
13 ^^Irm -f $(CUDA_EXECUTABLES)
```

```
1 #!/bin/bash
2 #PBS -q debug-g
3 #PBS -l select=1
4 #PBS -W group_list=gc64
5 #PBS -j oe
6
7 module purge
8 module load cuda
9
10 cd ${PBS_O_WORKDIR}
11 ./a.out 256
```

Contents of .cu file: kernel function

- "kernel" (sometimes "GPU kernel"): A function that runs on GPU
 - SYNTAX: An ordinary C/C++ function that returns nothing (void)
 - SYNTAX: Add `__global__` keyword beforehand

```
1 __global__ void f(...args...) { ...body... }  
2
```

Listing 1: kernel template

Contents of .cu file: launching kernel by a host

A host (CPU) launches a kernel to devices.

- Programmers must specify the number of threads by `<<nb, bs>>`
 - `nb`: Number of Blocks (per grid) (sometimes `gridDim`)
 - `bs`: Block Size (sometimes `blockDim`)
- `nb * bs` is the number of CUDA threads created

```
1      // ... code run on host ...
2
3      f<<gridDim, blockDim>>(...args...);
```

- `nb, bs` can be 1,2,3-Dimensional using type `dim3`

```
1      dim3 block(x_threads_block, y_threads_block);      // x, y(, z)
2      dim3 grid(x_blocks_grid, z_blocks_grid);           // x, y(, z)
3
4      f<<grid, block>>(...args...);
```

Interlude: register values programs can explicitly use

- Threads which executes a single instruction can be executed in parallel.
- So, programmers are expected to make the kernel visible to all threads as the same instruction.
- For this perspective, a unique ID of each thread (= the loop index) is the key.

In CUDA, each kernel can access its own thread ID through built-in variables on the Special Registers.

- `blockDim.{x,y}` = `bs` (the block size)
 - `gridDim.{x,y}` = `nb` (the number of blocks)
 - `threadIdx.{x,y}` = the thread ID within a block ($\in [0, bs)$)
 - `blockIdx.{x,y}` = the thread's block ID ($\in [0, nb)$)
- > `blockDim.x * blockIdx.x + threadIdx.x` could be the loop index

Sample program #2: hello_gpu.cu

```
1 #include <stdio.h>
2 #include <cuda_runtime.h>
3
4 __device__ void gpuAdd(int *number){ *number += 1; }
5 __global__ void callGpu(int *number){ gpuAdd(number); }
6
7 int main(){
8     int device_id = 0; cudaSetDevice(device_id); //device set up
9     int *a = (int*)malloc(sizeof(int)); *a = 0; //allocate memory on host(cpu)
10    int *a_dev = 0; cudaMalloc((void**)&a_dev, sizeof(int)); // allocate memory on gpu
11    cudaMemcpy(a_dev, a, sizeof(int), cudaMemcpyHostToDevice); // memcpy host -> device
12    // execute
13    callGpu<<<1, 1>>>(a_dev);
14    cudaDeviceSynchronize(); // Wait until GPU processing finishes.
15
16    cudaMemcpy(a, a_dev, sizeof(int), cudaMemcpyDeviceToHost); // memcpy device -> host
17    cudaFree(a_dev); // free
18
19    printf("ans: %d \n", *a); return 0; // display the answer
20 }
```

Data communication between H & Ds: overview

- Host memory and device memory are (basically) *separate*
- The device(D) cannot access data on the host(H) and vice versa by hardware
 - Access to another memory (including that of another device) causes `Segfault`
- Software need to explicitly specify when and what to communicate between H & D

Data communication between H & Ds: template to send

1. allocate data of the same size both on host and device

```
1 size_t sz = sizeof(int)*len;
2 int *a = (int*)malloc(sz);
3 int *a_dev = 0;  cudaMalloc((void**)&a_dev, sz);
4 // (void**)&a_dev is the address of the pointer variable (a_dev)
5 // Head address of GPU global memory is written to a_dev
```

Data communication between H & Ds: template to send

1. allocate data of the same size both on host and device

```
1 size_t sz = sizeof(int)*len;
2 int *a = (int*)malloc(sz);
3 int *a_dev = 0;  cudaMalloc((void**)&a_dev, sz);
4 // (void**)&a_dev is the address of the pointer variable (a_dev)
5 // Head address of GPU global memory is written to a_dev
```

2. the host works on the kinda initialization of the host data

```
1 for ( ... ) { a[i] = ... } // on host, initialize input data of kernel
```

Data communication between H & Ds: template to send

1. allocate data of the same size both on host and device

```
1 size_t sz = sizeof(int)*len;
2 int *a = (int*)malloc(sz);
3 int *a_dev = 0;  cudaMalloc((void**)&a_dev, sz);
4 // (void**)&a_dev is the address of the pointer variable (a_dev)
5 // Head address of GPU global memory is written to a_dev
```

2. the host works on the kinda initialization of the host data

```
1 for ( ... ) { a[i] = ... } // on host, initialize input data of kernel
```

3. copy the data to the device

```
1 cudaMemcpy(a_dev, a, sz, cudaMemcpyHostToDevice); // target address,
               source address, size, keyword
```

Data communication between H & Ds: template to send

1. allocate data of the same size both on host and device

```
1 size_t sz = sizeof(int)*len;
2 int *a = (int*)malloc(sz);
3 int *a_dev = 0;  cudaMalloc((void**)&a_dev, sz);
4 // (void**)&a_dev is the address of the pointer variable (a_dev)
5 // Head address of GPU global memory is written to a_dev
```

2. the host works on the kinda initialization of the host data

```
1 for ( ... ) { a[i] = ... } // on host, initialize input data of kernel
```

3. copy the data to the device

```
1 cudaMemcpy(a_dev, a, sz, cudaMemcpyHostToDevice); // target address,
               source address, size, keyword
```

4. pass the device pointer to the kernel

```
1 f<<nb, bs>>(a_dev, ...);
```

Data communication between H & Ds: template to retrieve

1. allocate data of the same size both on host and device

```
1 size_t sz = sizeof(int)*len;
2 int *r = (int*)malloc(sz);           // host memory
3 int *r_dev = 0;  cudaMalloc((void**)&r_dev, sz); // device memory
```

Data communication between H & Ds: template to retrieve

1. allocate data of the same size both on host and device

```
1 size_t sz = sizeof(int)*len;
2 int *r = (int*)malloc(sz);           // host memory
3 int *r_dev = 0;  cudaMalloc((void**)&r_dev, sz); // device memory
```

2. pass the device pointer of receptacle to the kernel

```
1 f<<nb, bs>>(...,r_dev, ...); // args must include pointer(s) of input
    and output
```


Data communication between H & Ds: template to retrieve

1. allocate data of the same size both on host and device

```
1 size_t sz = sizeof(int)*len;
2 int *r = (int*)malloc(sz);           // host memory
3 int *r_dev = 0;  cudaMalloc((void**)&r_dev, sz); // device memory
```

2. pass the device pointer of receptacle to the kernel

```
1 f<<nb, bs>>(...,r_dev, ...); // args must include pointer(s) of input
    and output
```

3. copy the data to the device

```
1 cudaMemcpy(r, r_dev, sz, cudaMemcpyDeviceToHost); // target address,
    source address, size, keyword
```

Host-Device synchronization

- A kernel call and the host overlap.
- Multiple kernel calls are serialized on the GPU side, by default (Host basically cannot control)
 - **Grid** is an abstraction of a one time launch of the kernel.
 - Grid is managed using a data struct FIFO queue called **Stream**
 - If you want to execute multiple kernel calls concurrently, you must use multiple Streams or Devices.
- `cudaDeviceSynchronize()` is an API to wait for the kernel to finish

```
1 h0();  
2 g0<<...,>>();  
3 h1();  
4 g1<<...,>>();  
5 cudaDeviceSynchronize();  
6 h2();
```

- `g0` might overlap with `h1`
- `g0` and `g1` do not overlap because they are assigned to the same stream 0
- `h2s` does not overlap with anything because of `cudaDeviceSynchronize()`

(+) "Programmer-free" data communication between H & Ds

- Recent NVIDIA GPUs support [Unified Memory](#) that **eliminate** the need for

Sample Code: `vecadd`

(+) "Programmer-free" data communication between H & Ds

- Recent NVIDIA GPUs support **Unified Memory** that **eliminate** the need for
 - explicit data movement between host and device memory

Sample Code: `vecadd`

(+) "Programmer-free" data communication between H & Ds

- Recent NVIDIA GPUs support **Unified Memory** that **eliminate** the need for
 - explicit data movement between host and device memory
 - dual pointer management

Sample Code: `vecadd`

(+) "Programmer-free" data communication between H & Ds

- Recent NVIDIA GPUs support **Unified Memory** that **eliminate** the need for
 - explicit data movement between host and device memory
 - dual pointer management
- **cudaMallocManaged** is an API that hides the copy of data between the host and devices

Sample Code: vecadd

(+) "Programmer-free" data communication between H & Ds

- Recent NVIDIA GPUs support **Unified Memory** that **eliminate** the need for
 - explicit data movement between host and device memory
 - dual pointer management
- **cudaMallocManaged** is an API that hides the copy of data between the host and devices
 - The CPU and GPU page tables are linked at the unified virtual address, and a hardware page fault fires the moment a GPU/CPU accesses a page.

Sample Code: vecadd

(+) "Programmer-free" data communication between H & Ds

- Recent NVIDIA GPUs support **Unified Memory** that **eliminate** the need for
 - explicit data movement between host and device memory
 - dual pointer management
- **cudaMallocManaged** is an API that hides the copy of data between the host and devices
 - The CPU and GPU page tables are linked at the unified virtual address, and a hardware page fault fires the moment a GPU/CPU accesses a page.
 - The moment a GPU/CPU accesses a page, a hardware page fault fires and the page is moved on-demand.

Sample Code: vecadd

(+) "Programmer-free" data communication between H & Ds

- Recent NVIDIA GPUs support **Unified Memory** that **eliminate** the need for
 - explicit data movement between host and device memory
 - dual pointer management
- **cudaMallocManaged** is an API that hides the copy of data between the host and devices
 - The CPU and GPU page tables are linked at the unified virtual address, and a hardware page fault fires the moment a GPU/CPU accesses a page.
 - The moment a GPU/CPU accesses a page, a hardware page fault fires and the page is moved on-demand.
 - Coherency is ensured at the kernel synchronization point.

Sample Code: `vecadd`

Conventional vs Unified Memory

Purpose	Conventional (<code>cudaMalloc</code> + <code>cudaMemcpy</code>)	Unified Memory (<code>cudaMallocManaged</code>)
Memory Management	allocates separate buffers for host and device, and makes copy explicit	single pointer can be referenced from either CPU/GPU
copy	All buffer transfers each time <code>cudaMemcpy()</code> is called	automatic transfer per page (on demand)
address space	different values for CPU and GPU	Unified Virtual Address (UVA) -share same value
oversubscribe	impossible	GPU Can allocate more than the memory capacity and swap unused pages to the host
Optimization API	None	Manual tuning of placement with <code>cudaMemPrefetchAsync</code> , <code>cudaMemAdvise</code>

HOW CUDA works

1. What is CUDA?: Introduction
2. A CUDA program for beginners
3. How CUDA works
 - 3.1 Architecture of NVIDIA GPU
 - 3.2 Grid, block, thread; abstractions by CUDA
 - 3.3 Warp; Parallel Thread eXecution
 - 3.4 Stream: beyond Grid
 - 3.5 Memory Hierarchy in CUDA
 - 3.6 Resolving race condition on CUDA
4. Optimize CUDA program
5. Practice Problems

Architecture of NVIDIA GPU: GPU unit

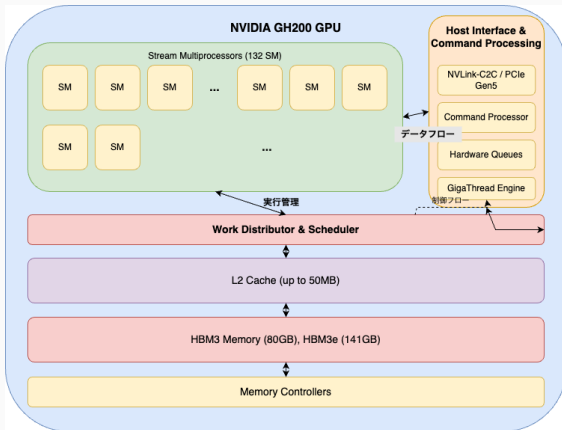
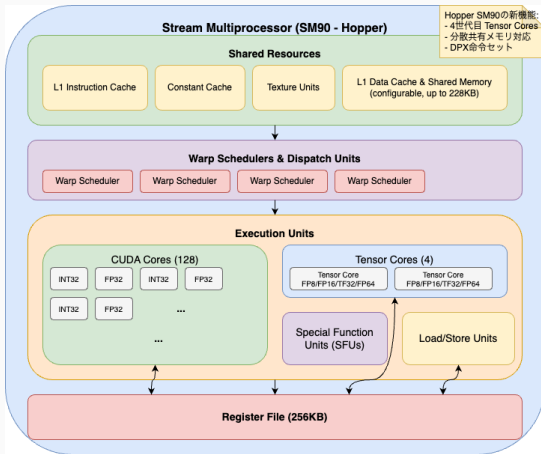


Figure 3: Device GPU

- **GBM3 Memory**
known as **global memory**, which has large capacity but slow access speed
- **Stream Multiprocessor (SM)**
In charge of multiple blocks, performs the operations that make up the kernel in parallel.
- **Host Interface & Command Processing**
Interface to communicate with the host CPU, and a command processing unit that manages the execution of commands.

Architecture of NVIDIA GPU: Stream Multiprocessors



- **shared memory**

has small capacity but relatively fast access speed

- **Warp scheduler**

Manage parallel execution on CUDA cores in units of Warps that comprise blocks allocated to the SM

- **CUDA core**: The core that performs the actual computation

- **Tensor core**: The core specialized for matrix operations

Figure 4: Device GPU

3 easy pieces about hardware - software

Programmers know 3 things about CUDA: Grid, Block, Thread

- (Review) **Grid** corresponds to a one time launch of the kernel.
 - A **Grid** is assigned to a single GPU unit (i.e., a single device)
 - (Review) **Grid** is a collection of **Blocks**
- A **Block** is the unit of dispatching to an SM
 - A **Block** is assigned to a single SM
i.e., once a block starts running, it stays on the SM (= occupies registers and shared memory) all the way until it finishes
 - (Review) A **Block** is a collection of **Threads**
- A **Thread** is the smallest unit of execution
- A **Thread** is assigned to a single CUDA core

Motivation: You may have questions like ...

1. Is it allowed for a `block` to have **more threads than the number of processors(CUDA cores)** in the allocated **SM**, and if so, how is this handled?
2. Is it allowed to have **more blocks than the number of SMs** that make up the **GPU unit** corresponding to the `grid`, and if so, how is this handled?
3. How is the execution across **multiple grids** handled?

Hints:

- `threads` belonging to the same `block` are executed by the same **SM**, this does **NOT** mean they are executed in parallel.
- A **SM** is assigned to a `block`, but this does **NOT** mean that an **SM** can only be in charge of one `block`.

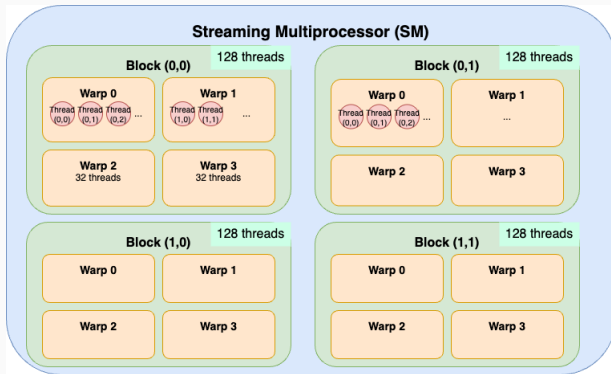
Warp: The way to realize "Parallel" Thread eXecution

- The unit of **instruction execution** in the SM is a **Warp**
- The number of threads that make up a warp has always been **32**.
- Each thread in a warp shares an instruction pointer (i.e, executes the same instruction at the same time).
- The warp scheduler selects a warp from the ready queue (Warp pool) and executes the instruction.

Hierarchy within an SM

Parallelism within an Stream Multiprocessor consists of three levels.

$\text{thread} \subset \text{warp} \subset \text{block} \subset \text{SM}$



- (recap) A group of 32 CUDA threads makes a warp
- A group of $bs/32$ warps makes a block
- There are multiple blocks active on a single SM

Limitation of Hardware: performance degradation

Stream

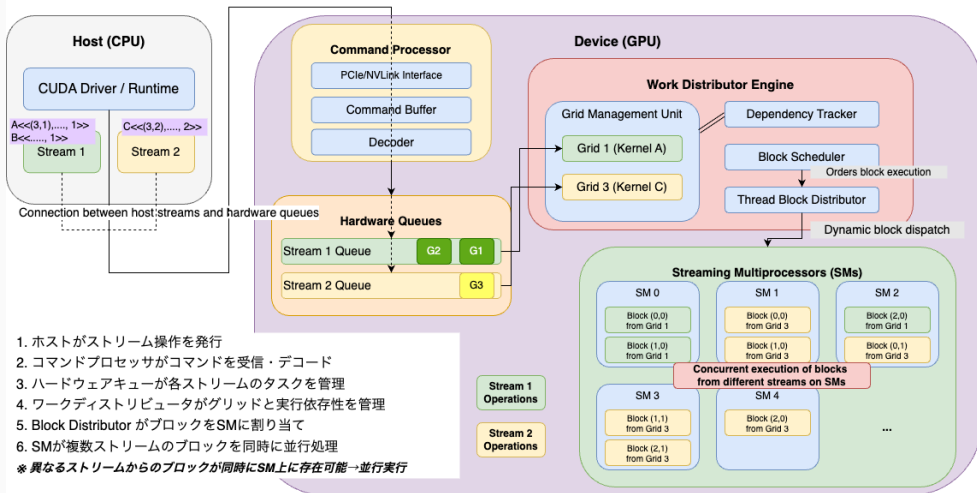
(review) As long as multiple kernels are submitted to the same stream(i.e., the default stream), they are always executed in series on the GPU side.

∴ If program doesn't specify Stream when launching a kernel, the grid is automatically assigned to **legacy stream 0**.

- **Stream** is a **FIFO queue** that binds the sequence of operations passed by the host to the GPU.
*The host runtime writes the operation to the command buffer and queues up entries with the same **stream ID** ; the GPU grabs the queue on **doorbell notification** and distributes it to the hardware engine (Device), keeping each stream in order.*
- **Grid** (= a task) is assigned to a Stream.
- If program assign grids to multiple Streams (= launch kernels), GPU firmware pop an item from the **available** queue with the **highest priority**
- Scheduling policy between Streams is a complete **blackbox**.

Flow of Stream Operation

ストリーム操作の流れと処理



Interlude: template to use multiple Streams

```
1  cudaStream_t sCompute, sCopy;
2  cudaStreamCreate(&sCompute);
3  cudaStreamCreate(&sCopy);
4
5  // 1) 非同期コピー (→) をHD stream sCopy
6  cudaMemcpyAsync(d_in, h_in, bytes, cudaMemcpyHostToDevice, sCopy);
7
8  // 2) カーネルを stream sCompute
9  myKernel<<<grid, block, 0, sCompute>>>(d_in, d_out); // 0: sharedMem ID,
10
11 // 3) 非同期コピー (→) をDH stream sCopy
12 cudaMemcpyAsync(h_out, d_out, bytes, cudaMemcpyDeviceToHost, sCopy);
13
14 // 4) 任意の同期点
15 cudaEvent_t done; cudaEventCreate(&done);
16 cudaEventRecord(done, sCopy); // sCopy 完了後に立つ
17 cudaStreamWaitEvent(sCompute, done, 0); // sCompute は done まで待つ
```

(+) Using multiple Devices

Memory hierarchy of NVIDIA GPU

Tiling: effective use of shared memory

To avoid race conditions: effective use of parameters

To avoid race conditions: barrier synchronization

To avoid race conditions: Cooperative groups

(Reluctantly) resolve race conditions: Atomic accumulations

1. What is CUDA?: Introduction
2. A CUDA program for beginners
3. How CUDA works
4. Optimize CUDA program
5. Practice Problems
 - 5.1 k-NN for vector Database
 - 5.2 SW's local alignment; DP acceleration with GPU

Further problems ...